

Course #4:

Recurrent Neural Networks

Roadmap

- Recap from course #3
- Auto-encoders
- Recurrent Neural Networks

Lecture. #3

Things to know (AE)

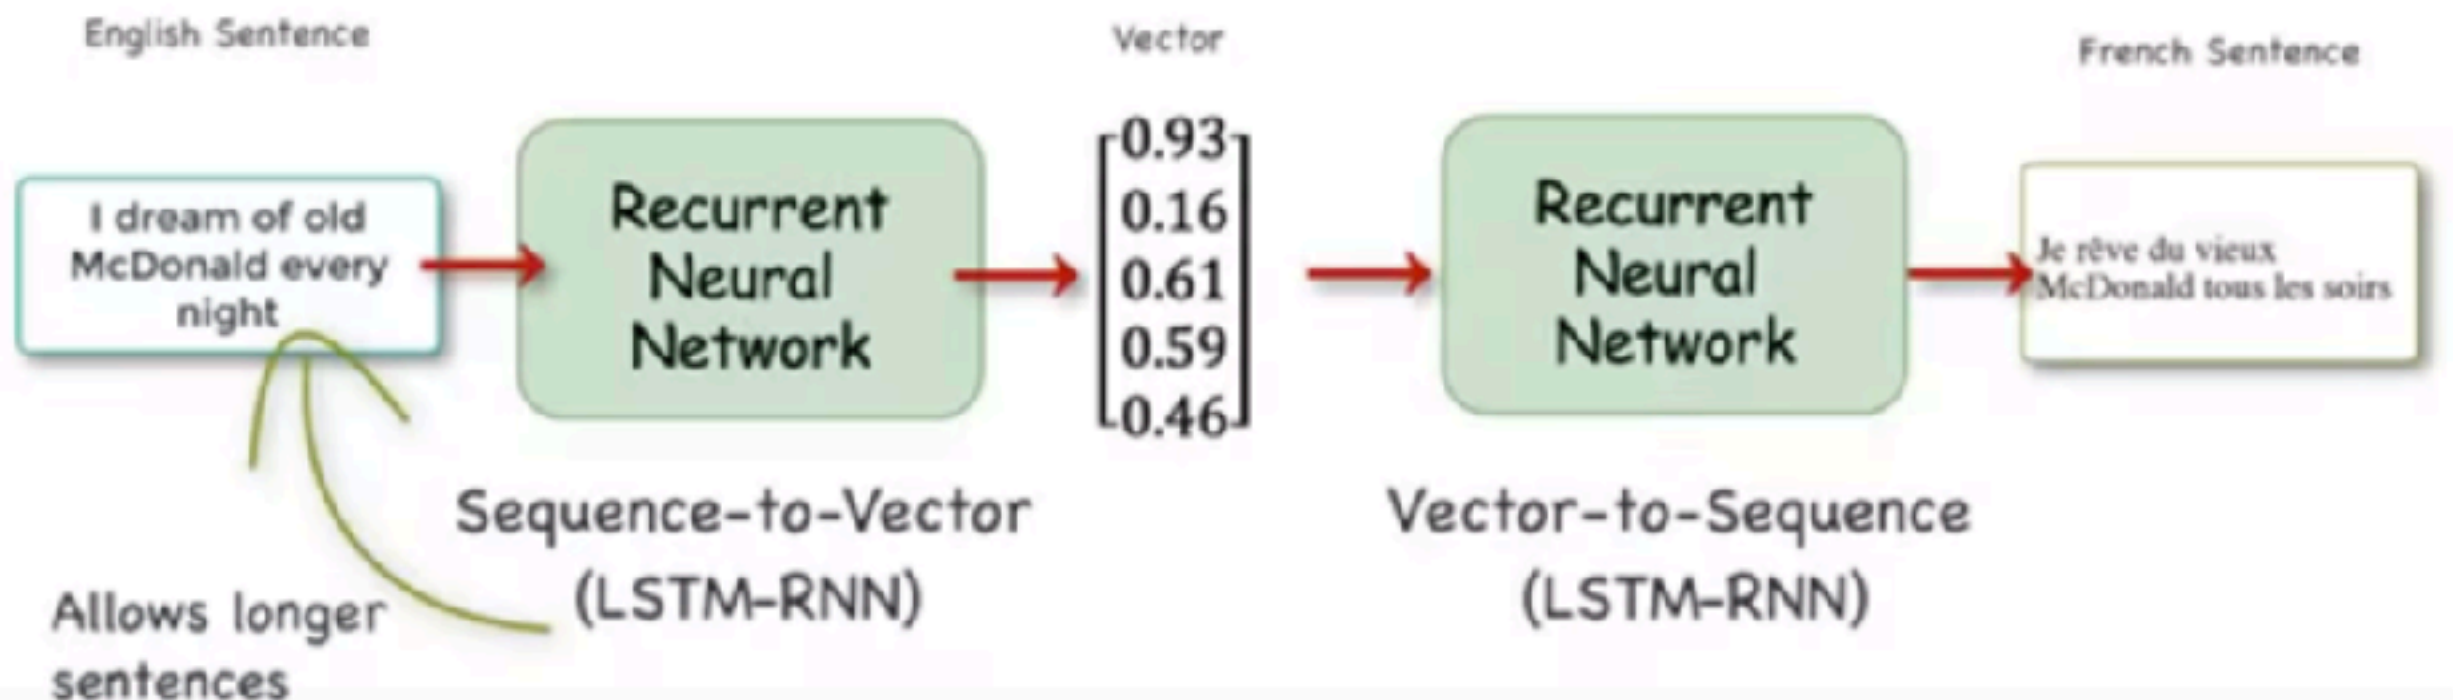
- Auto-encoder
- Latent variable
- UNet
- ResNet

Recurrent Neural Networks

Application to text data

Language Translation

Encoder-Decoder Architecture



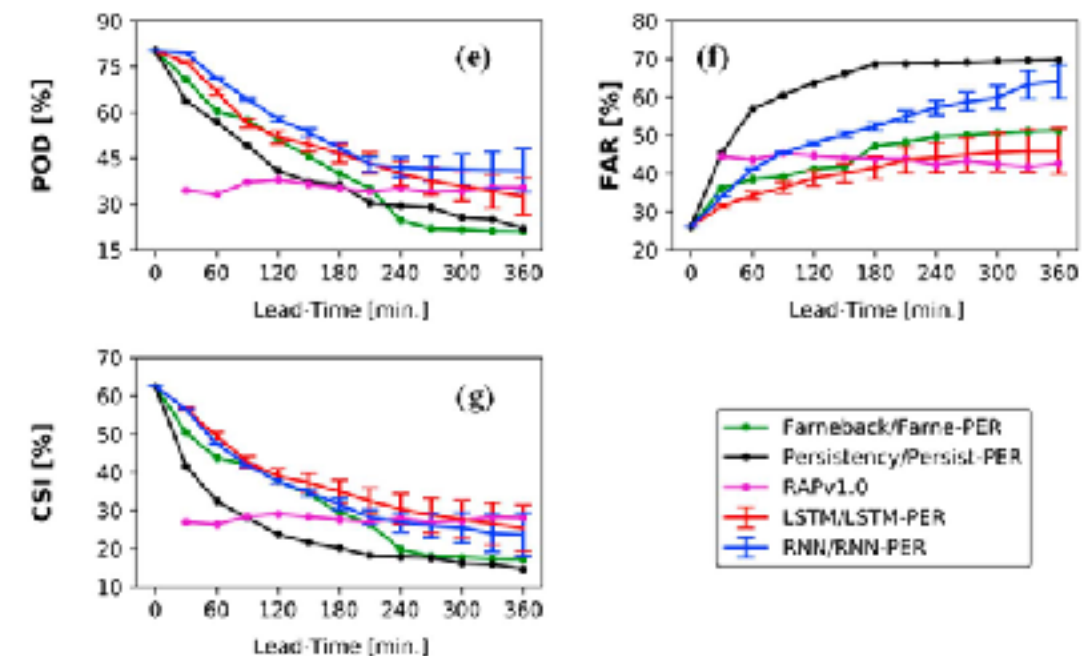

Applications to geoscience

IEEE JOURNAL OF SELECTED TOPICS IN APPLIED EARTH OBSERVATIONS AND REMOTE SENSING, VOL. 13, 2020

2853

A Deep Learning Method With Merged LSTM Neural Networks for SSHA Prediction

Tao Song¹, Senior Member, IEEE, Jingyu Jiang, Wei Li, and Danya Xu



AGU100 ADVANCING EARTH AND SPACE SCIENCE

Journal of Geophysical Research: Atmospheres

RESEARCH ARTICLE

10.1029/2018JD028375

Key Points:


- Artificial intelligence techniques are useful tools in support of forecasting complex precipitation in short range (0–6 hr).
- Long Short-Term Memory structure is capable of learning spatial and temporal correlations efficiently.
- The framework provides accurate precipitation forecasts especially for

Short-Term Precipitation Forecast Based on the PERSIANN System and LSTM Recurrent Neural Networks

Ata Akbari Asanjan¹, Tiantian Yang^{1,2}, Kuolin Hsu^{1,2}, Soroosh Sorooshian¹, Junqiang Lin⁴, and Qidong Peng⁴

¹Department of Civil and Environmental Engineering, Center for Hydrometeorology and Remote Sensing, University of California, Irvine, CA, USA, ²School of Civil Engineering and Environmental Science, University of Oklahoma, Norman, OK, USA, ³Center for Excellence for Ocean Engineering, National Taiwan Ocean University, Keelung, Taiwan, ⁴China Institute of Water Resources and Hydropower Research, Beijing, China

Recurrent Neural Networks



Applications:

- Time-series forecasting
- Audio processing
- Translation
-

Do CNNs also apply ?


RNN: Underlying formulation (similar to state-space representation)

$x \rightarrow$ input vector

$y \rightarrow$ output vector


$h \rightarrow$ hidden state

$t \rightarrow$ index (time, index, etc...)



$$h_t = f(W_{hh}h_{t-1} + W_{xh}x_t) \text{ and } y_t = g(W_{hy}h_t)$$


RNN: Underlying formulation (similar to state-space representation)



Classic RNN Architectures


- **Conventional RNN with sigmoid**

- The sensitivity of the input values decays over time
- The network forgets the previous input



- **Long-Short Term Memory (LSTM)** [2]

- The cell remember the input as long as it wants
- The output can be used anytime it wants



Dense and convolutional versions of LSTM and GRU exist
(depending on the structure of the hidden state)

Short-term forecasting application (L63 case-study)

https://github.com/CIA-Oceanix/DLOA2023/blob/main/lectures/notebooks/notebook_PytorchLightning_Forecasting_L63_students.ipynb

Lecture. #3

Things to know (RNN)

- Recurrent Neural Network
- LSTM
- Unfloded and folded representations

Physics-informed/theory-guided networks

General question

How to exploit physical knowledge in the design of neural networks ?

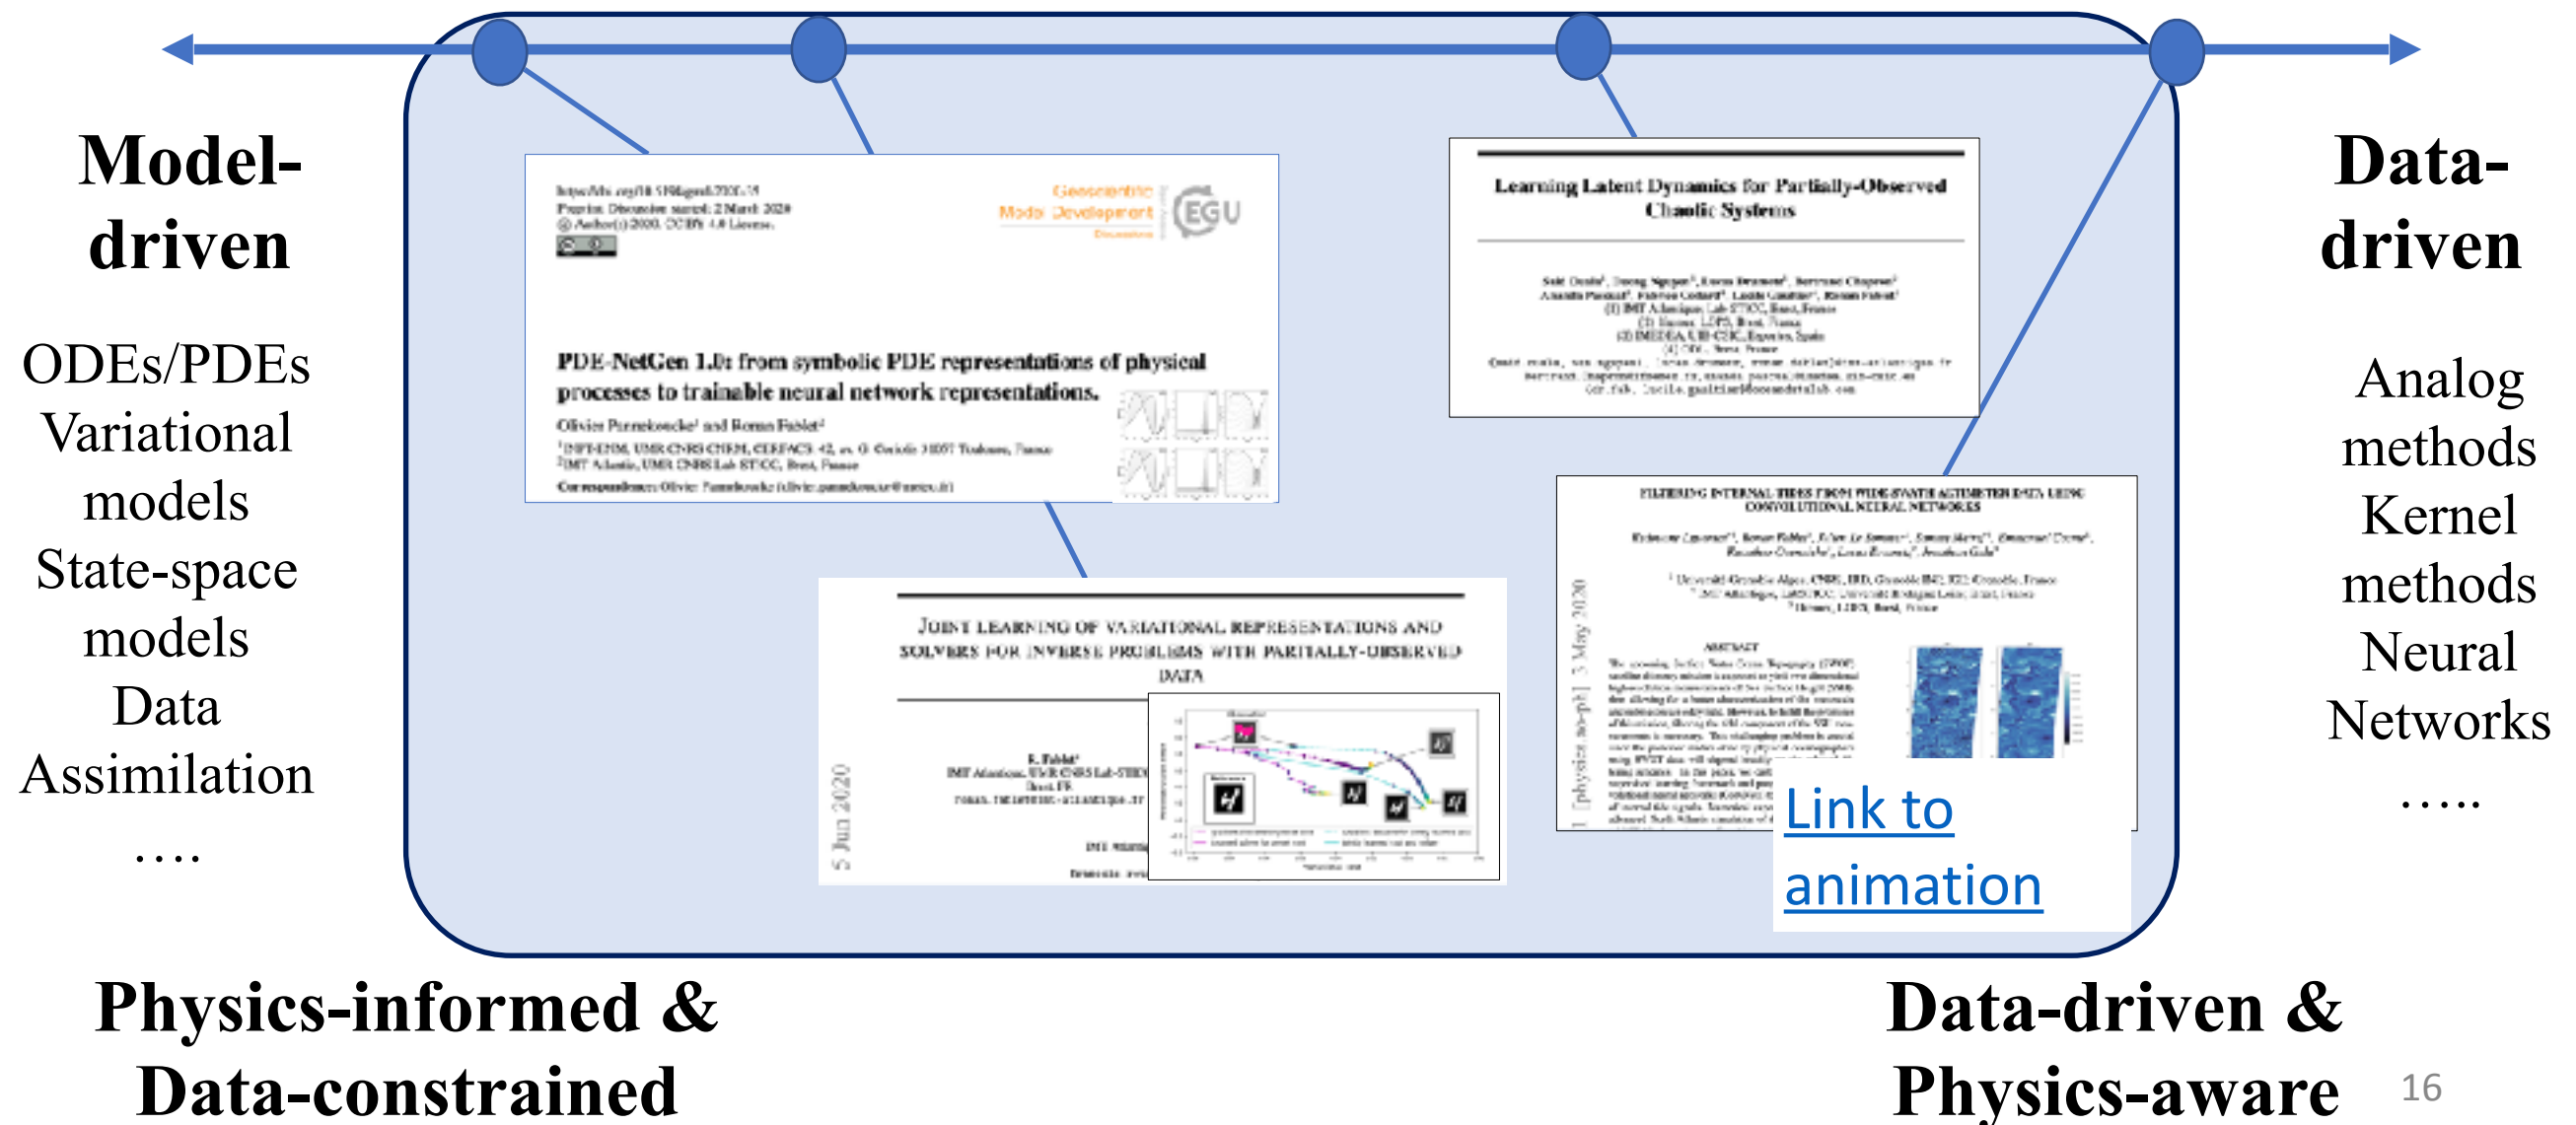
Bridging physics & AI: a broader picture



Making the most of AI and Physics Theory

- Model-Driven/Theory-Guided & Data-Constrained schemes
- Data-Driven & Physics-Aware schemes (eg, Ouala et al., 2019)

Bridging Physics & AI: a broader picture



Bridging physics & AI: a broader picture



Making the most of AI and Physics Theory


- **Model-Driven/Theory-Guided & Data-Constrained schemes**
- Data-Driven & Physics-Aware schemes (eg, Ouala et al., 2019)

How to embed physics-driven priors in DL models ?

An illustration through L63 dynamics: numerical experiments (Fablet et al., 2018)

$$\begin{aligned}\frac{dx(t)}{dt} &= \sigma (y(t) - x(t)) \\ \frac{dy(t)}{dt} &= x(t) (\rho - z(t)) - y(t) \\ \frac{dz(t)}{dt} &= x(t) y(t) - \beta z(t)\end{aligned}$$


Lorenz-63 equations



Associated Euler integration scheme
 $d_t X(t) = F_\theta (X(t))$


➡ $X(t + \delta) = X(t) + \delta \cdot F_\theta (X(t))$

NN architecture for differential operator



Bilinear architecture

NN architecture for integration scheme



ResNet architecture (Residual Network)

How to embed physics-driven priors in DL models ?


An illustration through L63 dynamics: numerical experiments (Fablet et al., 2018)

$$\begin{aligned}\frac{dx(t)}{dt} &= \sigma (y(t) - x(t)) \\ \frac{dy(t)}{dt} &= x(t) (\rho - z(t)) - y(t) \\ \frac{dz(t)}{dt} &= x(t) y(t) - \beta z(t)\end{aligned}$$

Lorenz-63 equations


Generalization to higher-order integration schemes (eg, RK4)

$$d_t X(t) = F_\theta (X(t))$$



$$X(t + \delta) = X(t) + \sum_i \beta_i k_i$$

with $k_i = F_\theta (X(t) + \delta \alpha_i k_{i-1})$



NB: Same number of trainable model parameters as the Euler-based architecture

How to embed physics-driven priors in DL models ?

Implementation for L63 dynamics

https://github.com/CIA-Oceanix/DLOA2023/blob/main/lectures/notebooks/notebook_PytorchLightning_Forecasting_L63_students.ipynb

How to embed physics-driven priors in DL models ?

An illustration through L63 dynamics: numerical experiments (Fablet et al., 2018)


Forecasting experiments

Noise-free training data

Forecasting time step	t_0+h	t_0+4h	t_0+8h
Analog forecasting	$<10^{-6}$	0.002	0.005
Sparse regression	$<10^{-6}$	0.002	0.006
MLP	$<10^{-6}$	0.018	0.044
<i>Bi-NN(4)</i>	$<10^{-6}$	$<10^{-6}$	$<10^{-6}$

Noisy training data ($\sigma=0.5$)

Forecasting time step	t_0+h	t_0+4h	t_0+8h
Analog forecasting	$<10^{-6}$	2.01	2.2
<i>Bi-NN(4)</i>	$<10^{-6}$	0.054	0.14



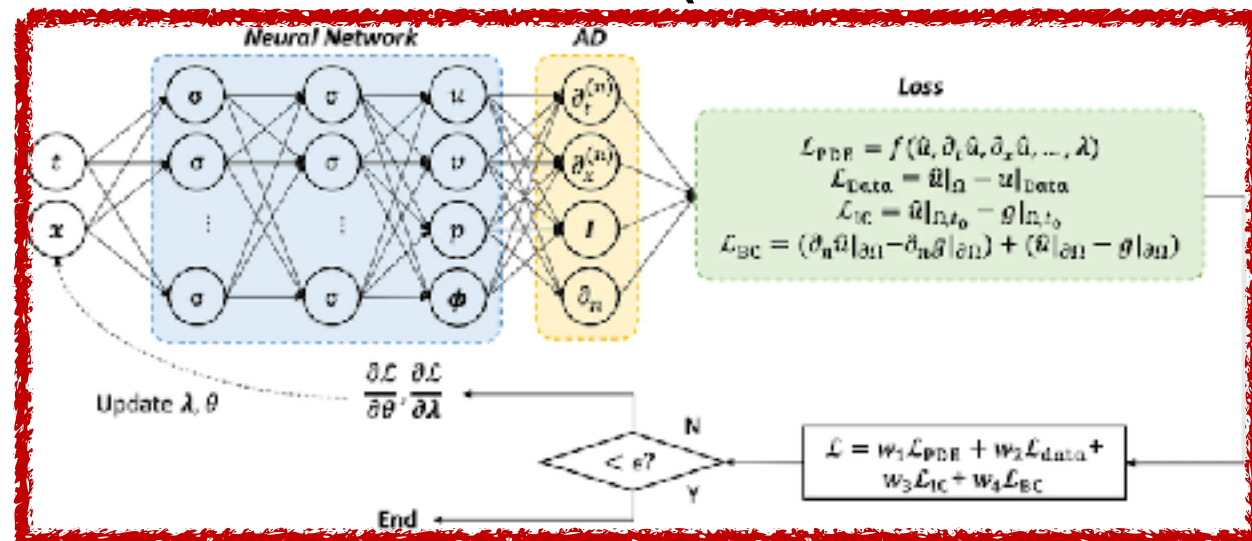
Assimilation experiment

(1 obs. every 8 time steps)

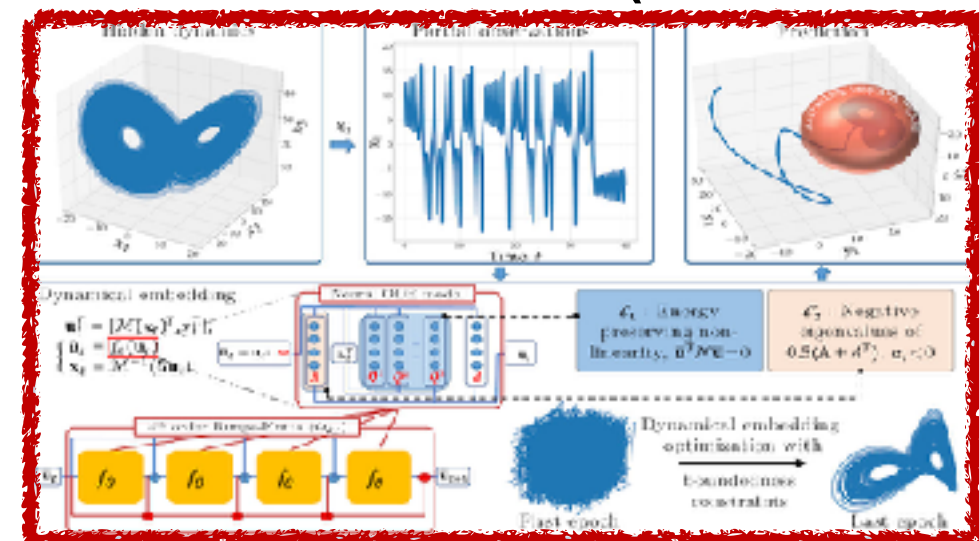
Noise standard deviation in training data	0	0.25	1
<u>True model</u>	<u>0.50</u>	-	-
Analog forecasting	0.65	1.17	1.81
<i>Bi-NN(4)</i>	0.60	0.75	0.86

Physics-informed neural architectures

PINNS for ODEs/PDEs (from Cai et al., 2021)



Identification of ODEs (Ouala et al., 2021)



Other physics-informed representation: Hamiltonian representations [eg, Greydanus et al., 2019], Fourier representations [eg, Li et al., 2021]

Variational data assimilation (Fablet et al., 2021)

From a Variational DA formulation

$$\hat{x} = \arg \min_x \|y - H(x)\|^2 + \lambda \|x - \Phi(x)\|^2$$

Trainable variational model

Trainable gradient-based solver

Associated end-to-end scheme

