



# Pytorch Lightning

Quentin Febvre, PhD student



## “Simple” training routine walkthrough

```
1 n_epochs = 30
2 model_1 = MLP()
3 model_1.to(device=device)
4 optimizer = torch.optim.SGD(model_1.parameters(), lr = 0.01)
5 criterion = nn.CrossEntropyLoss()
6
7 train_losses, valid_losses = [], []
8 valid_loss_min = np.Inf
9
10 for epoch in range(n_epochs):
11     train_loss, valid_loss = 0, 0
12
13     model.train()
14     for data, label in train_loader:
15         data = data.to(device=device, dtype=torch.float32)
16         label = label.to(device=device, dtype=torch.long)
17         optimizer.zero_grad()
18         output = model(data)
19         loss = criterion(output, label)
20         loss.backward()
21         optimizer.step()
22         train_loss += loss.item() * data.size(0)
23
24     model.eval()
25     for data, label in valid_loader:
26         data = data.to(device=device, dtype=torch.float32)
27         label = label.to(device=device, dtype=torch.long)
28         with torch.no_grad():
29             output = model(data)
30             loss = criterion(output, label)
31             valid_loss += loss.item() * data.size(0)
32
33     train_loss /= len(train_loader.sampler)
34     valid_loss /= len(valid_loader.sampler)
35     train_losses.append(train_loss)
36     valid_losses.append(valid_loss)
37
38     print('epoch: {} \ttraining Loss: {:.6f} \tvalidation Loss: {:.6f}'.format(
39         epoch+1, train_loss, valid_loss))
40
41     if valid_loss <= valid_loss_min:
42         print('validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
43             valid_loss_min,
44             valid_loss))
45         torch.save(model.state_dict(), 'model.pt')
46         valid_loss_min = valid_loss
```

```

1 n_epochs = 30
2 model_1 = MLP()
3 model_1.to(device=device)
4 optimizer = torch.optim.SGD(model_1.parameters(), lr = 0.01)
5 criterion = nn.CrossEntropyLoss()
6
7 train_losses, valid_losses = [], []
8 valid_loss_min = np.Inf
9
10 for epoch in range(n_epochs):
11     train_loss, valid_loss = 0, 0
12
13     model.train()
14     for data, label in train_loader:
15         data = data.to(device=device, dtype=torch.float32)
16         label = label.to(device=device, dtype=torch.long)
17         optimizer.zero_grad()
18         output = model(data)
19         loss = criterion(output, label)
20         loss.backward()
21         optimizer.step()
22         train_loss += loss.item() * data.size(0)
23
24     model.eval()
25     for data, label in valid_loader:
26         data = data.to(device=device, dtype=torch.float32)
27         label = label.to(device=device, dtype=torch.long)
28         with torch.no_grad():
29             output = model(data)
30             loss = criterion(output, label)
31         valid_loss += loss.item() * data.size(0)
32
33     train_loss /= len(train_loader.sampler)
34     valid_loss /= len(valid_loader.sampler)
35     train_losses.append(train_loss)
36     valid_losses.append(valid_loss)
37
38     print('epoch: {} \ttraining Loss: {:.6f} \tvalidation Loss: {:.6f}'.format(
39         epoch+1, train_loss, valid_loss))
40
41     if valid_loss <= valid_loss_min:
42         print('validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
43             valid_loss_min,
44             valid_loss))
45         torch.save(model.state_dict(), 'model.pt')
46         valid_loss_min = valid_loss

```

“Simple” learning routine contains many considerations intertwined:

- Training logic:
  - What do I optimize for ?
  - How do I compute it ?
  - How do I optimize it ?

```

1  n_epochs = 30
2  model_1 = MLP()
3  model_1.to(device=device)
4  optimizer = torch.optim.SGD(model_1.parameters(), lr = 0.01)
5  criterion = nn.CrossEntropyLoss()
6
7  train_losses, valid_losses = [], []
8  valid_loss_min = np.Inf
9
10 for epoch in range(n_epochs):
11     train_loss, valid_loss = 0, 0
12
13     model.train()
14     for data, label in train_loader:
15         data = data.to(device=device, dtype=torch.float32)
16         label = label.to(device=device, dtype=torch.long)
17         optimizer.zero_grad()
18         output = model(data)
19         loss = criterion(output, label)
20         loss.backward()
21         optimizer.step()
22         train_loss += loss.item() * data.size(0)
23
24     model.eval()
25     for data, label in valid_loader:
26         data = data.to(device=device, dtype=torch.float32)
27         label = label.to(device=device, dtype=torch.long)
28         with torch.no_grad():
29             output = model(data)
30             loss = criterion(output, label)
31             valid_loss += loss.item() * data.size(0)
32
33     train_loss /= len(train_loader.sampler)
34     valid_loss /= len(valid_loader.sampler)
35     train_losses.append(train_loss)
36     valid_losses.append(valid_loss)
37
38     print('epoch: {} \ttraining Loss: {:.6f} \tvalidation Loss: {:.6f}'.format(
39         epoch+1, train_loss, valid_loss))
40
41     if valid_loss <= valid_loss_min:
42         print('validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
43             valid_loss_min,
44             valid_loss))
45         torch.save(model.state_dict(), 'model.pt')
46         valid_loss_min = valid_loss

```

“Simple” learning routine contains many considerations intertwined:

- Training logic:
  - What do I optimize for ?
  - How do I compute it ?
  - How do I optimize it ?
- Training flow:
  - Loop for each epoch
  - Iterate over each minibatch
- Autograd/Pytorch flow
  - train/eval mode
  - Backprop + parameter update
  - No\_grad (disable grad computation)

```

1 n_epochs = 30
2 model = MLP()
3 model.to(device=device)
4 optimizer = torch.optim.SGD(model.parameters(), lr = 0.01)
5 criterion = nn.CrossEntropyLoss()
6
7 train_losses, valid_losses = [], []
8 valid_loss_min = np.Inf
9
10 for epoch in range(n_epochs):
11     train_loss, valid_loss = 0, 0
12
13     model.train()
14     for data, label in train_loader:
15         data = data.to(device=device, dtype=torch.float32)
16         label = label.to(device=device, dtype=torch.long)
17         optimizer.zero_grad()
18         output = model(data)
19         loss = criterion(output, label)
20         loss.backward()
21         optimizer.step()
22         train_loss += loss.item() * data.size(0)
23
24     model.eval()
25     for data, label in valid_loader:
26         data = data.to(device=device, dtype=torch.float32)
27         label = label.to(device=device, dtype=torch.long)
28         with torch.no_grad():
29             output = model(data)
30             loss = criterion(output, label)
31             valid_loss += loss.item() * data.size(0)
32
33     train_loss /= len(train_loader.sampler)
34     valid_loss /= len(valid_loader.sampler)
35     train_losses.append(train_loss)
36     valid_losses.append(valid_loss)
37
38     print('epoch: {} \ttraining Loss: {:.6f} \tvalidation Loss: {:.6f}'.format(
39         epoch+1, train_loss, valid_loss))
40
41     if valid_loss <= valid_loss_min:
42         print('validation loss decreased {:.6f} --> {:.6f}. Saving model ...'.format(
43             valid_loss_min,
44             valid_loss))
45         torch.save(model.state_dict(), 'model.pt')
46         valid_loss_min = valid_loss

```

“Simple” learning routine contains many considerations intertwined:

- Training logic:
  - What do I optimize for ?
  - How do I compute it ?
  - How do I optimize it ?
- Training flow:
  - Loop for each epoch
  - Iterate over each minibatch
- Autograd/Pytorch flow
  - train/eval mode
  - Backprop + parameter update
  - No\_grad (disable grad computation)
- Engineering flow / Plugins
  - Training hardware (.to(device))
  - Logging training metrics
  - Saving best model



```

1 n_epochs = 30
2 model 1 = MLP()
3 model 1.to(device=device)
4 optimizer = torch.optim.SGD(model 1.parameters(),lr = 0.01)
5 criterion = nn.CrossEntropyLoss()
6
7 train_losses, valid_losses = [], []
8 valid_loss_min = np.Inf
9
10 for epoch in range(n_epochs):
11     train_loss, valid_loss = 0, 0
12
13     model.train()
14     for data, label in train_loader:
15         data = data.to(device=device, dtype=torch.float32)
16         label = label.to(device=device, dtype=torch.long)
17         optimizer.zero_grad()
18         output = model(data)
19         loss = criterion(output, label)
20         loss.backward()
21         optimizer.step()
22         train_loss += loss.item() * data.size(0)
23
24     model.eval()
25     for data, label in valid_loader:
26         data = data.to(device=device, dtype=torch.float32)
27         label = label.to(device=device, dtype=torch.long)
28         with torch.no_grad():
29             output = model(data)
30             loss = criterion(output, label)
31         valid_loss += loss.item() * data.size(0)
32
33     train_loss /= len(train_loader.sampler)
34     valid_loss /= len(valid_loader.sampler)
35     train_losses.append(train_loss)
36     valid_losses.append(valid_loss)
37
38     print('epoch: {} \ttraining Loss: {:.6f} \tvalidation Loss: {:.6f}'.format(
39         epoch+1, train_loss, valid_loss))
40
41     if valid_loss <= valid_loss_min:
42         print('validation loss decreased {:.6f} --> {:.6f}). Saving model ...'.format(
43             valid_loss_min,
44             valid_loss))
45         torch.save(model.state_dict(), 'model.pt')
46         valid_loss_min = valid_loss

```

“Simple” learning routine contains many considerations intertwined:

#### - Training logic:

- What do I optimize for ?
- How do I compute it ?
- How do I optimize it ?

Extract

#### - Training flow:

- Loop for each epoch
- Iterate over each minibatch

#### - Autograd/Pytorch flow

- train/eval mode
- Backprop + parameter update
- No\_grad (disable grad computation)

Abstract

#### - Engineering flow / Plugins

- Training hardware (.to(device))
- Logging training/validation metrics
- Saving best model

Configure

```

1 n_epochs = 30
2 model = MLP()
3 model.to(device=device)
4 optimizer = torch.optim.SGD(model.parameters(), lr = 0.01)
5 criterion = nn.CrossEntropyLoss()

6
7 train_losses, valid_losses = [], []
8 valid_loss_min = np.Inf
9
10 for epoch in range(n_epochs):
11     train_loss, valid_loss = 0, 0
12
13     model.train()
14     for data, label in train_loader:
15         data = data.to(device=device, dtype=torch.float32)
16         label = label.to(device=device, dtype=torch.long)
17         optimizer.zero_grad()
18         output = model(data)
19         loss = criterion(output, label)
20         loss.backward()
21         optimizer.step()
22         train_loss += loss.item() * data.size(0)
23
24     model.eval()
25     for data, label in valid_loader:
26         data = data.to(device=device, dtype=torch.float32)
27         label = label.to(device=device, dtype=torch.long)
28         with torch.no_grad():
29             output = model(data)
30             loss = criterion(output, label)
31             valid_loss += loss.item() * data.size(0)
32
33     train_loss /= len(train_loader.sampler)
34     valid_loss /= len(valid_loader.sampler)
35     train_losses.append(train_loss)
36     valid_losses.append(valid_loss)
37
38     print('epoch: {} \ttraining Loss: {:.6f} \tvalidation Loss: {:.6f}'.format(
39         epoch+1, train_loss, valid_loss))
40
41     if valid_loss <= valid_loss_min:
42         print('validation loss decreased {:.6f} --> {:.6f}). Saving model ...'.format(
43             valid_loss_min,
44             valid_loss))
45         torch.save(model.state_dict(), 'model.pt')
46         valid_loss_min = valid_loss

```

```

4 class LitMLP(pl.LightningModule):
5     def __init__(self): ...
18
19     def forward(self, x): ...
27
28     def training_step(self, batch, batch_idx):
29         data, label = batch
30         output = self(data)
31         return F.cross_entropy(output, label)
32
33     def validation_step(self, batch, batch_idx):
34         data, label = batch
35         output = self(data)
36         return F.cross_entropy(output, label)
37
38     def configure_optimizers(self):
39         return torch.optim.SGD(self.parameters(), lr = 0.01)
40

```

## Training logic extraction:

- Change the nn.Module to pytorch\_lightning.LightningModule
- Compute and return the optimization objective in the training and validation step method

```

1 n_epochs = 30
2 model = MLP()
3 model.to(device=device)
4 optimizer = torch.optim.SGD(model.parameters(), lr = 0.01)
5 criterion = nn.CrossEntropyLoss()
6
7 train_losses, valid_losses = [], []
8 valid_loss_min = np.Inf
9
10 for epoch in range(n_epochs):
11     train_loss, valid_loss = 0, 0
12
13     model.train()
14     for data, label in train_loader:
15         data = data.to(device=device, dtype=torch.float32)
16         label = label.to(device=device, dtype=torch.long)
17         optimizer.zero_grad()
18         output = model(data)
19         loss = criterion(output, label)
20         loss.backward()
21         optimizer.step()
22         train_loss += loss.item() * data.size(0)
23
24     model.eval()
25     for data, label in valid_loader:
26         data = data.to(device=device, dtype=torch.float32)
27         label = label.to(device=device, dtype=torch.long)
28         with torch.no_grad():
29             output = model(data)
30             loss = criterion(output, label)
31             valid_loss += loss.item() * data.size(0)
32
33     train_loss /= len(train_loader.sampler)
34     valid_loss /= len(valid_loader.sampler)
35     train_losses.append(train_loss)
36     valid_losses.append(valid_loss)
37
38     print('epoch: {} \ttraining Loss: {:.6f} \tvalidation Loss: {:.6f}'.format(
39         epoch+1, train_loss, valid_loss))
40
41     if valid_loss <= valid_loss_min:
42         print('validation loss decreased {:.6f} --> {:.6f}). Saving model ...'.format(
43             valid_loss_min,
44             valid_loss))
45         torch.save(model.state_dict(), 'model.pt')
46         valid_loss_min = valid_loss

```

```

4 class LitMLP(pl.LightningModule):
5     def __init__(self): ...
18
19     def forward(self, x): ...
27
28     def training_step(self, batch, batch_idx):
29         data, label = batch
30         output = self(data)
31         return F.cross_entropy(output, label)
32
33     def validation_step(self, batch, batch_idx):
34         data, label = batch
35         output = self(data)
36         return F.cross_entropy(output, label)
37
38     def configure_optimizers(self):
39         return torch.optim.SGD(self.parameters(), lr = 0.01)
40
1 model = LitMLP()
2 trainer = pl.Trainer(max_epochs=1)
3 trainer.fit(
4     model, train_dataloaders=train_loader, val_dataloaders=valid_loader)

```

## Training flow abstraction:

- Instantiate the `pytorch_lightning.Trainer` object
- Pass the Lightning module and the training data to the `fit` method to perform the optimization procedure



```

1 n_epochs = 30
2 model = MLP()
3 model.to(device=device)
4 optimizer = torch.optim.SGD(model.parameters(), lr = 0.01)
5 criterion = nn.CrossEntropyLoss()
6
7 train_losses, valid_losses = [], []
8 valid_loss_min = np.Inf
9
10 for epoch in range(n_epochs):
11     train_loss, valid_loss = 0, 0
12
13     model.train()
14     for data, label in train_loader:
15         data = data.to(device=device, dtype=torch.float32)
16         label = label.to(device=device, dtype=torch.long)
17         optimizer.zero_grad()
18         output = model(data)
19         loss = criterion(output, label)
20         loss.backward()
21         optimizer.step()
22         train_loss += loss.item() * data.size(0)
23
24     model.eval()
25     for data, label in valid_loader:
26         data = data.to(device=device, dtype=torch.float32)
27         label = label.to(device=device, dtype=torch.long)
28         with torch.no_grad():
29             output = model(data)
30             loss = criterion(output, label)
31             valid_loss += loss.item() * data.size(0)
32
33     train_loss /= len(train_loader.sampler)
34     valid_loss /= len(valid_loader.sampler)
35     train_losses.append(train_loss)
36     valid_losses.append(valid_loss)
37
38     print('epoch: {} \ttraining Loss: {:.6f} \tvalidation Loss: {:.6f}'.format(
39         epoch+1, train_loss, valid_loss))
40
41     if valid_loss <= valid_loss_min:
42         print('validation loss decreased {:.6f} --> {:.6f}. Saving model ...'.format(
43             valid_loss_min,
44             valid_loss))
45         torch.save(model.state_dict(), 'model.pt')
46         valid_loss_min = valid_loss

```

```

4 class LitMLP(pl.LightningModule):
5     def __init__(self): ...
18
19     def forward(self, x): ...
27
28     def training_step(self, batch, batch_idx):
29         data, label = batch
30         output = self(data)
31         loss = F.cross_entropy(output, label)
32         self.log('train loss', loss)
33         return loss
34
35     def validation_step(self, batch, batch_idx):
36         data, label = batch
37         output = self(data)
38         loss = F.cross_entropy(output, label)
39         self.log('val loss', loss)
40         return loss
41
42     def configure_optimizers(self):
43         return torch.optim.SGD(self.parameters(), lr = 0.01)
44

```

```

1 model = LitMLP()
2 model_checkpoint = pl.callbacks.ModelCheckpoint(monitor='val_loss')
3 logger = pl.loggers.CSVLogger('logs', name='mlp_mnist')
4 trainer = pl.Trainer(
5     max_epochs=10,
6     gpus=1,
7     callbacks=[model_checkpoint],
8     logger=logger
9 )
10
11 trainer.fit(model, train_dataloaders=train_loader, val_dataloaders=valid_loader)
12

```

## Engineering/Plugins configuration:

- Instantiate the logger and model checkpoint callback
- Pass as parameter to the trainer, as well as other flags (gpus...)

## Vanilla Pytorch:



VS

## Pytorch Lightning:



- + All the logic is exposed sequentially: easier to write and reason about
- + Fewer classes/concept to handle
- Code harder to read
- Reinvent the wheel each time

- + Code is organized : easier to read/share
- + Less code
- + Benefit from SOTA implementations
- Additional concepts to understand: Trainer, Logger, ModelCheckpointCallback
- Some modifications are less straight-forward (find the method, the parameter...)

Give it a try :

*[notebooks/notebook\\_PytorchLightning\\_MNIST\\_CNN\\_students.ipynb](#)*